

Abstract

We study efficient serving of long-context language models with sparse attention and a chunked KV-cache scheduler that keeps memory bounded. The scheduler preserves answer quality on long-document benchmarks.

1 Introduction

Serving long contexts is expensive in memory and compute.

References

- [1] Vaswani et al. Attention Is All You Need. NeurIPS 2017.
- [2] He et al. Deep Residual Learning for Image Recognition. CVPR 2015.
- [3] Chen et al. Chunked KV-Cache Scheduling for Long Contexts. MLSys 2024.