

Abstract

We fine-tune language models with 4-bit quantized LoRA adapters under a strict memory budget using a budget-aware rank scheduler.

1 Introduction

Adapters reduce memory but quantization adds representation error.

References

- [1] Vaswani et al. Attention Is All You Need. NeurIPS 2017.
- [2] He et al. Deep Residual Learning for Image Recognition. CVPR 2015.
- [3] Dettmers et al. QLoRA Efficient Finetuning of Quantized LLMs. NeurIPS 2023.